

ECHOES: Equal-weight Completed Hypothetical Observation Ensembles for BOSS CMASS-South: Survey-Ready Posterior Samples for Arbitrary Galaxy-Clustering Statistics

Tom Abel

June 18, 2026

Abstract

Most precision galaxy-clustering analyses communicate between survey teams and theory teams through a small set of summary statistics, most often two-point correlation functions or power spectra. This convention is powerful because the observational effects that enter these statistics have been studied for decades, but it also makes each new statistic responsible for reinterpreting the survey window, fiber assignment, redshift failures, imaging systematics, and random catalogs. We present ECHOES, Equal-weight Completed Hypothetical Observation Ensembles, for BOSS DR12 CMASS-South: an ensemble of equal-weight, cosmology-free completed catalogs that keep every observed galaxy at its measured (α, δ, z) and restore the spectroscopically missing galaxies at their real SDSS imaging positions. The unknown redshifts are sampled from a data-driven posterior combining a color-space k -nearest-neighbor photo- z , an empirical close-pair prior, and a local-density line-of-sight prior. Smooth imaging-systematic weights are represented by local analog galaxies rather than by per-object completeness weights. The released posterior is compact: one 2.04 MB file stores the fixed observed catalog and the inverse CDF for each restored target, and a numpy-only sampler generates reproducible complete catalogs from integer seeds. For seed 0 the sample contains 119,923 equal-weight galaxies, of which 109,636 are observed, 5,272 are restored fiber-collision targets, 1,505 are restored redshift failures, and 3,510 are imaging-systematic analogs. We validate the product against truth-known completions and MultiDark-Patchy mocks, finding that projected clustering is recovered to the 1–2% level over $0.5 < r_p < 40 h^{-1}\text{Mpc}$, while the posterior completion uncertainty in $w_p(r_p)$ is only 0.17–0.44% per bin, below the sample-variance scale of the BOSS South Galactic Cap. We also describe a fully complementary graphGP route: a conditional anisotropic Gaussian-process posterior over the galaxy density field, sampled by Matheron’s rule on a sparse graph and evaluated along each missing galaxy’s line of sight. The resulting catalogs are intended to let analysts run two-point, higher-order, marked, nearest-neighbor, topological, field-level,

or simulation-based statistics directly on survey-ready data while propagating the observational-completion uncertainty by resampling the catalog posterior.

Keywords: Large-scale structure of the universe (902); Galaxy surveys (598); Astrostatistics (1882); Cosmological parameters (339); Observational cosmology (1146)

1 Introduction

The central compression in observational large-scale structure is not a simulation or an estimator, but an agreement about what object is being communicated. A survey team measures a summary statistic from a complicated instrumental and astrophysical selection function; a theory team predicts the same statistic for a cosmological and galaxy-formation model. The summary statistic is the interface. For modern spectroscopic galaxy surveys that interface has usually been a two-point statistic: a correlation function, a projected correlation function, a redshift-space multipole, or a Fourier-space power spectrum [Davis and Peebles, 1983, Landy and Szalay, 1993, Hamilton, 1993, Feldman et al., 1994]. This choice is not merely historical. The observational corrections for two-point clustering have been revisited repeatedly, calibrated in mocks, and encoded in survey catalogs, random catalogs, and weights.

At the same time, the cosmological information in a galaxy survey is not exhausted by two-point clustering. The literature contains a long sequence of proposed and demonstrated summaries: higher-order moments and hierarchical statistics [Szapudi and Szalay, 1993, Fry and Gaztanaga, 1993], void-probability functions [Einasto et al., 1991], counts-in-cells and one-point density PDFs [Baugh et al., 1995, Szapudi et al., 1996, Uhlemann et al., 2020, Friedrich et al., 2020], marked correlation functions and marked power spectra [Harker et al., 2006, Satpathy et al., 2019, Massara et al., 2021, 2023], nearest-neighbor distributions [Banerjee and Abel, 2021, Yuan et al., 2024], wavelet and scattering-transform statistics [Cheng et al., 2020, Valogiannis and Dvorkin, 2022, Regalado-Saint Blancard et al., 2024], persistent-homology summaries of the cosmic web [Wilding et al., 2021, Biagetti et al., 2022, Yip et al., 2024], and field-level

or simulation-based inference methods [Jasche and Wandelt, 2013, Nguyen et al., 2024, Lemos et al., 2024, Hahn et al., 2023a,b, 2024b,a, Massara et al., 2025]. Large simulation suites such as Quijote [Villaescusa-Navarro et al., 2020] have been essential for calibrating and stress-testing many of these summaries. The Beyond-2pt Collaboration’s parameter-masked mock-data challenge [Beyond-2pt Collaboration et al., 2025] makes the same point in a deliberately broad form: once the Gaussian, two-point description is relaxed, many summaries can extract complementary information, and their relative utility is an empirical question.

The barrier between those ideas and real spectroscopic data is often not the statistic itself. It is the observational model. A statistic that is easy to define on a periodic simulation box can become difficult to interpret in a survey with angular holes, variable targeting, fiber collisions, redshift failures, stellar density, seeing, extinction, and redshift-dependent selection. For two-point statistics, decades of work have made the corresponding correction pathways usable. For many newer statistics, the same corrections are either rederived from scratch, approximated, or avoided by restricting the analysis to mocks and forecasts. This paper asks whether the interface can be moved one step earlier: instead of releasing only weighted galaxies and randoms and asking every statistic to understand the weights, release a posterior over completed galaxy catalogs that have already absorbed the main observational incompleteness into equal-weight point sets.

This perspective has a long historical lineage. The first large maps of extragalactic structure already had to confront the fact that the sky is not a uniform observing surface. Hubble [1934] discussed the distribution of extragalactic nebulae in the presence of Galactic obscuration, a problem that later became a quantitative dust-correction program culminating in all-sky reddening maps such as Schlegel et al. [1998]. Redshift surveys then turned angular inhomogeneity into a three-dimensional selection problem [Davis and Peebles, 1983, de Lapparent et al., 1986]. The 2dF Galaxy Redshift Survey and SDSS established the modern industrial mode of galaxy redshift surveys [York et al., 2000, Colless et al., 2001, Cole et al., 2005], and the SDSS luminous red galaxy sample made the two-point correlation function a precision cosmological observable with the detection of the baryon acoustic feature [Eisenstein et al., 2005].

BOSS inherited this history and made it operational. Its spectrographs, fiber system, target selection, and large-scale-structure catalogs are documented in Dawson et al. [2013], Smee et al. [2013], Reid et al. [2016]. The final DR12 cosmology analyses used a mature system of angular masks, radial random catalogs, redshift-failure weights, close-pair weights, imaging-systematic weights, and FKP weights [Alam et al., 2017, Ross et al., 2017]. The resulting measurements were robust because the effect of these ingredients on the two-point statistics had been studied. But the same survey artifacts have no rea-

son to couple only to two-point functions. Fiber collisions remove close angular pairs; redshift failures correlate with observing conditions and target properties; imaging systematics modulate the target density across degree scales; and the random catalog carries a particular representation of the angular and radial selection function.

The same issues will not become simpler in later surveys. DESI’s robotic fiber-positioner system and its fiber-assignment algorithm create a richer assignment problem than the BOSS plug-plate geometry [Silber et al., 2023, Pinol et al., 2017, Lan et al., 2023, Lasker et al., 2025]. Euclid, the Nancy Grace Roman Space Telescope, the Chinese Space Station Telescope, and the Vera C. Rubin Observatory’s LSST will produce enormous galaxy samples with their own mixtures of imaging, spectroscopy, selection, and calibration systematics [Euclid Collaboration et al., 2025, Wang et al., 2022, Miao et al., 2023, CSST Collaboration et al., 2026, Zhan and Tyson, 2018]. If those surveys wish to serve a community that will analyze the data with statistics not yet standardized, the most useful data product may not be a single corrected summary statistic. It may be a survey-ready probabilistic catalog.

This work builds such a product for BOSS DR12 CMASS-South. We call the approach ECHOES: Equal-weight Completed Hypothetical Observation Ensembles, echoes of the survey one would have taken if the principal spectroscopic incompleteness and smooth imaging systematics were negligible. It is not a new cosmological measurement. It is a data-modeling and validation paper. We construct an ensemble of completed catalogs conditioned on the published BOSS data and SDSS imaging: observed spectroscopic galaxies remain fixed, missing targets are restored at real imaging positions, and redshift uncertainty is sampled. The ensemble is cosmology-free in the sense that it contains only RA, Dec, and redshift; an analyst chooses a fiducial cosmology only when measuring three-dimensional clustering. The ECHOES product is designed so that the summary statistic, rather than the observational correction, can once again be the object under study.

2 ECHOES Completed-Survey Methodology

2.1 Data model and scope

The target ECHOES product is an ensemble

$$\mathcal{C}^{(s)} = \{\alpha_i^{(s)}, \delta_i^{(s)}, z_i^{(s)}, p_i^{(s)}\}_{i=1}^{N_s}, \quad (1)$$

where s is a seed and p_i is a provenance code. The coordinates are observed RA, Dec, and redshift. No comoving coordinate, Alcock–Paczynski scaling [Alcock and Paczynski, 1979], redshift-space model [Kaiser, 1987], or fiducial cosmology is stored in the catalog. The completion posterior is conditional on the observed BOSS cat-

alog, the SDSS imaging target catalog, and the BOSS angular mask and weights. It does not represent cosmic variance; it represents the uncertainty in replacing spectroscopic incompleteness and smooth imaging-systematic weights by equal-weight galaxies.

The ECHOES product corrects three effects in the default release: fiber collisions, redshift failures, and the smooth imaging-systematic modulation represented by `WEIGHT_SYSTOT`. An optional inpainted version also fills interior mask holes, but the default scientific product leaves the survey mask explicit. The FKP weight [Feldman et al., 1994] is not a completeness correction and is not converted into galaxies. Analysts may still apply FKP or other inverse-variance weights when estimating a statistic; the released point catalog itself is equal weight.

2.2 BOSS CMASS-South sample

We use the BOSS DR12 CMASS South Galactic Cap subset with the same clean cuts used in several modern analyses of these data:

$$\text{RA} < 28^\circ \text{ or } \text{RA} > 335^\circ, \quad \text{Dec} > -6^\circ, \quad 0.45 < z < 0.60 \quad (2)$$

The input contains 109,636 observed spectroscopic galaxies. The BOSS completeness weights imply a mean completeness correction of 1.074, or about 7.4% missing objects before the smooth imaging-systematic analog step. In the real target recovery used here, 5,272 fiber-collision targets and 1,505 redshift-failure targets are restored, for 6,777 spectroscopically missing galaxies in the compact posterior package.

The BOSS loader reads the standard large-scale-structure columns `RA`, `DEC`, `Z`, `WEIGHT_SYSTOT`, `WEIGHT_NOZ`, `WEIGHT_CP`, and `WEIGHT_FKP`. It also reads `MODELFLUX`, `EXTINCTION`, `IMATCH`, and `ICOLLIDED` when photometry is requested. Extinction-corrected Pogson magnitudes are computed as

$$m_b = 22.5 - 2.5 \log_{10} f_b - A_b, \quad (3)$$

with adjacent colors $(u - g, g - r, r - i, i - z)$. The photometric-redshift model uses $(g - r, r - i, i - z, i)$, dropping the u band because it is noisy for the red CMASS population.

2.3 Recovering the missing target list

The completion uses real imaging positions for missing galaxies. We start from the DR12 CMASS color-selected target catalog, cross-matched to the good redshift BOSS LSS catalog. A target with a spectrum but nonzero `ZWARNING` is classified as a redshift failure. A target with no usable spectrum is eligible to be a fiber-collision target. Because the no-spectrum photometric pool also contains objects that were not part of the final BOSS LSS sample, it is tied back to the BOSS book-keeping. For each observed survivor with `WEIGHT_CP`

1, the algorithm greedily claims $\text{round}(\text{WEIGHT_CP} - 1)$ nearest unclaimed no-spectrum targets within the BOSS fiber-collision scale, $62''$. Redshift failures are tied to survivors with `WEIGHT_NOZ` > 1 using the analogous $\text{round}(\text{WEIGHT_NOZ} - 1)$ count and a wider $15'$ neighborhood. This preserves the integer count implied by the official weights while placing the restored galaxies at their actual imaging positions and colors.

2.4 Color-space photo- z posterior

For each missing target we require a redshift posterior, not only a point estimate. We fit a nonparametric k -nearest-neighbor model to the observed CMASS galaxies in the whitened four-dimensional feature space $(g - r, r - i, i - z, i)$. The default model uses $k = 100$ and a `ckdTree`. For a query feature vector $\mathbf{x}$, the posterior support is the set of neighbor redshifts $\{z_j\}_{j=1}^k$ with inverse-distance weights

$$w_j(\mathbf{x}) = \frac{(d_j^2 + \epsilon)^{-1}}{\sum_{\ell=1}^k (d_\ell^2 + \epsilon)^{-1}}. \quad (4)$$

The held-out diagnostic in the report gives $\sigma_{\text{NMAD}} = 0.019$, a bias of -0.0006 , a catastrophic outlier rate of 1.8%, and a probability-integral-transform mean of 0.503. These numbers are adequate for the role the photo- z plays here: it supplies a color likelihood that is then multiplied by a local-density prior along the same line of sight.

2.5 Local-density redshift posterior

The default redshift engine is a fast local-density approximation to a conditional field posterior. For missing target i at angular position $\hat{\mathbf{n}}_i$ and colors $\mathbf{x}_i$, redshifts are sampled from

$$p_i(z) \propto \hat{\lambda}_{\text{loc}}(\hat{\mathbf{n}}_i, z) \hat{p}_{\text{col}}(z | \mathbf{x}_i) q_i(z), \quad (5)$$

where $\hat{p}_{\text{col}}$ is the color-space photo- z density, q_i is an empirical close-pair prior for collisions and unity for redshift failures, and $\hat{\lambda}_{\text{loc}}$ is the observed local line-of-sight density. The local density is estimated from the $K = 150$ nearest observed spectroscopic galaxies in angular unit-vector space:

$$\hat{\lambda}_{\text{loc}}(\hat{\mathbf{n}}_i, z) = \sum_{j \in \mathcal{N}_K(i)} \exp \left[-\frac{(z - z_j)^2}{2\sigma_f^2} \right], \quad \sigma_f = 0.004. \quad (6)$$

The photo- z factor is evaluated on the same 256-point redshift grid with a Gaussian smoothing scale $\sigma_p = 0.02$ applied to the neighbor redshifts and their k NN weights. For collided targets, $q_i(z)$ is the symmetrized histogram of redshift differences among observed close pairs within $62''$, measured directly from the data over $|\Delta z| < 0.06$. The sampled redshift is then jittered by a Gaussian of width $\sigma_f/2$ and clipped to the CMASS redshift interval.

Equation (5) is deliberately observational. It does not assign a missing galaxy to a simulated halo or require a




Figure 1: BOSS CMASS-South geometry and completeness inputs. The completion is conditioned on the observed BOSS footprint, the published completeness weights, and the SDSS imaging positions of spectroscopically missing targets.




Figure 2: Color-space k NN photo- z diagnostic for the CMASS sample. The photo- z is used as a likelihood factor for missing redshifts, not as a standalone redshift assignment.

fiducial cosmology. The object is placed at its real angular position and on a redshift structure seen in the neighboring spectroscopy, with the color likelihood preventing the local-density prior from assigning obviously incompatible redshifts.

2.6 Imaging-systematic analogs and provenance

The BOSS imaging-systematic correction `WEIGHT_SYSTOT` is a smooth density correction, not a statement that a particular nearby target failed to receive a fiber. We therefore do not convert it into exact duplicate galaxies. For every observed or spectroscopically restored base object, the sampler adds

$$n_{\text{extra}} = \lfloor \max(w_{\text{systot}} - 1, 0) + U \rfloor, \quad U \sim \mathcal{U}(0, 1), \quad (7)$$

local analogs at the same source position with a $1''$ Gaussian angular jitter and the same redshift. This preserves resolved angular and three-dimensional clustering while avoiding a zero-separation spike that would corrupt nearest-neighbor or counts-in-cells statistics. The default implementation keeps all real detections; it restores density deficits implied by $w_{\text{systot}} > 1$ rather than thinning observed galaxies in regions with $w_{\text{systot}} < 1$.

Each output object carries a provenance code: 0 for observed spectroscopy, 1 for restored fiber-collision target, 2 for restored redshift failure, 3 for imaging-systematic analog, 4 for rare host-redshift fallback, and 5 for optional mask inpainting. The provenance vector is central to reproducibility because analysts can rerun a statistic

on individual components or test sensitivity to the imaging analog prescription.

2.7 Compact posterior and sampling

The default posterior has a simple product structure once the observed data are fixed. The observed galaxies never change across realizations. The missing galaxy positions never change. The redshift posterior for each missing galaxy, Equation (5), is also fixed after it is computed. A realization is only a draw from each per-object inverse CDF plus the stochastic imaging analogs.

The release stores the posterior as a compressed `npz` file. It contains the observed redshifts, the base RA and Dec arrays for observed plus restored targets, the base provenance vector, `WEIGHT_SYSTOT`, and an inverse CDF for each of the 6,777 missing targets. The inverse CDF is tabulated at 65 quantile levels and quantized to `uint16`. The file `cmass_south_posterior.npz` is 2.04 MB. A matching uniform-footprint random catalog, `cmass_south_randoms.npz`, contains 438,544 points. Because CMASS-South is approximately 99% complete in the angular random mask, these randoms are uniform over the footprint to the percent level.

Sampling is deterministic under a seed:

1. draw one uniform variate for each missing target and invert its stored redshift CDF;
2. concatenate the fixed observed galaxies and the restored targets;
3. draw local analog counts from `WEIGHT_SYSTOT`;
4. return arrays `ra`, `dec`, `z`, `prov`, and `N`.

For seed 0, the sampled catalog contains 119,923 objects: 109,636 observed, 5,272 collision restorations, 1,505 redshift-failure restorations, and 3,510 smooth-systematic analogs. Without the imaging analogs, the fixed base catalog contains 116,413 objects.

2.8 Clustering measurements and acceptance metrics

For angular clustering we measure Landy-Szalay estimates [Landy and Szalay, 1993] with the fixed ran-




Figure 3: Missing-target recovery and sampled redshift distributions. The restored galaxies are placed at real imaging positions, and only their redshifts are sampled. Realizations differ in the redshifts of missing galaxies and in the stochastic imaging-systematic analogs.

dom catalog. For three-dimensional statistics we use **Corrfunc** [Sinha and Garrison, 2020]. The catalog itself is cosmology-free; **Corrfunc** measurements convert redshift to comoving distance at measurement time, using the analyst’s fiducial cosmology, and then call the OpenMP mock pair counters with `is_comoving_dist=True`. The main validation statistics are $w(\theta)$, $w_p(r_p)$, $\xi(r_p, \pi)$, and redshift-space multipoles $\xi_0(s)$ and $\xi_2(s)$. We also evaluate higher-order or non-two-point summaries: nearest-neighbor CDFs and counts-in-cells. These are deliberately included because the purpose of the catalog posterior is to support statistics beyond those for which BOSS weights were optimized.

The acceptance criteria are empirical. In truth-known tests, the completed ensemble should recover the truth catalog to the percent level over the scales where the input mocks themselves are faithful. On the real data, the equal-weight completed catalogs should reproduce the official $w_c = \text{WEIGHT_SYSTOT}(\text{WEIGHT_CP} + \text{WEIGHT_NOZ} - 1)$ weighted BOSS clustering within the known scale of the correction. The completion posterior spread should be smaller than the sample-variance covariance for BOSS CMASS-South, because it represents only uncertainty in observational completion conditioned on the observed catalog.

3 The graphGP Complement

The local-density engine in Section 2.5 is fast and compressible, but it treats the missing-galaxy redshift posteriors as conditionally independent once the observed catalog is fixed. The graphGP route is a complementary, more explicit field-level model. It assumes that the complete galaxy overdensity field can be approximated by a Gaussian process with covariance tied to the observed two-point function, and that the observed galaxies are noisy samples of that field.

Let $\delta(\mathbf{x})$ be the galaxy overdensity in a fiducial comov-

ing coordinate system. The graphGP prior is

$$\delta(\mathbf{x}) \sim \mathcal{GP}[0, K_\xi(|\mathbf{x} - \mathbf{x}'|)], \quad (8)$$

where $K_\xi(r)$ is a tabulated covariance derived from a Landy–Szalay measurement of the galaxy correlation function. The likelihood is a Gaussian approximation to an inhomogeneous Poisson process. At observed galaxy locations $\mathbf{x}_D$, the data vector is

$$y_i = \bar{n}_i^{-1} - 1, \quad N_{ii} = \bar{n}_i^{-1}, \quad (9)$$

where $\bar{n}_i$ is the local mean density implied by the survey selection. The posterior is Gaussian:

$$\mu_{\text{post}}(\mathbf{x}) = K(\mathbf{x}, \mathbf{x}_D) [K_{DD} + N]^{-1} y_D, \quad (10)$$

$$C_{\text{post}}(\mathbf{x}, \mathbf{x}') = K(\mathbf{x}, \mathbf{x}') - K(\mathbf{x}, \mathbf{x}_D) [K_{DD} + N]^{-1} K(\mathbf{x}_D, \mathbf{x}'). \quad (11)$$

The implementation samples this posterior by Matheron’s rule [Rybicki and Press, 1992]. A prior draw δ_{prior} is generated on a sparse neighbor graph using a Vecchia-style approximation [Katzfuss et al., 2020, Schäfer et al., 2021]. The conditional draw is

$$\delta_{\text{post}}(\mathbf{x}) = \delta_{\text{prior}}(\mathbf{x}) + K(\mathbf{x}, \mathbf{x}_D) [K_{DD} + N]^{-1} [y_D - \delta_{\text{prior}}(\mathbf{x}_D)]. \quad (12)$$

The expensive solve is amortized over many missing-galaxy redshift draws. The primary stored object is a light-cone-native array $1 + \delta_{\text{post}}$ on redshift shells and HealPIX pixels. For a missing galaxy, graphGP replaces the local KDE in Equation (5) by the corresponding field draw:

$$p_i(z | s) \propto [1 + \delta_{\text{post}}^{(s)}(\hat{\mathbf{n}}_i, z)] \bar{n}(z) \hat{p}_{\text{col}}(z | \mathbf{x}_i) q_i(z). \quad (13)$$

All missing galaxies in realization s share the same density-field draw, so their redshift assignments are correlated.

This model has a fiducial comoving coordinate system internally because a three-dimensional covariance needs a distance. That coordinate system is a modeling gauge

for the field interpolation, not a cosmology baked into the released KNN product. The report tests this explicitly by rebuilding the graphGP route with different fiducial distance-redshift relations. Planck-like and Einstein-de Sitter choices give per-object redshift RMS differences of about 0.001–0.0012, correlations of 0.9995–0.9996, and changes in $w_p(r_p)$ below 0.1%. Within the accuracy demanded by the present BOSS completion, graphGP is therefore best viewed as a correlated posterior engine whose output is stable to reasonable internal coordinate choices.

4 Application to BOSS DR12 CMASS-South

4.1 Why BOSS is the right test case

BOSS DR12 is the ideal laboratory because the official two-point analyses are mature. The survey, target selection, and LSS catalogs are documented [Dawson et al., 2013, Reid et al., 2016], the spectroscopic system is understood [Smei et al., 2013], the DR12 cosmology and systematic-correction papers provide community reference measurements [Alam et al., 2017, Ross et al., 2017], and independent mock suites such as MultiDark-Patchy exist for validation [Kitauro et al., 2016]. A completed-catalog product can therefore be judged against a well-established reference: it should reproduce the weighted BOSS two-point measurements while exposing a point catalog suitable for statistics that were not part of the official BOSS pipeline.

The official BOSS correction strategy keeps the observed spectroscopic catalog and supplies per-object weights. Close-pair and redshift-failure weights move the missing object’s statistical weight to an observed neighbor, while `WEIGHT_SYSTOT` corrects smooth density modulations associated with imaging and observing conditions. Random catalogs encode the angular mask and radial selection. This system is extremely effective for the published two-point program, but for a general statistic the analyst must decide how to apply weights, whether the statistic is linear or nonlinear in those weights, and how random catalogs should enter. Fiber-collision corrections illustrate the point: nearest-neighbor upweighting and more refined pairwise or angular correction schemes can be calibrated for two-point functions [Guo et al., 2012, Mohammad et al., 2020], but they do not automatically define how a nonlinear statistic should respond to a missing close pair.

SIMBIG is a useful modern comparison point. It uses BOSS CMASS-South-like cuts and a forward-modeling program to connect nonlinear galaxy clustering statistics to cosmological parameters [Hahn et al., 2023a,b, 2024b,a]. Later work extends the same style of analysis to summaries such as the wavelet scattering transform and marked power spectra [Regalado-Saint Blancard

et al., 2024, Massara et al., 2025]. The scientific ambition is close to ours—to use information beyond the Gaussian two-point limit—but the observational interface remains the published BOSS selection and weighting system. Our product is complementary: it provides an ensemble of completed data catalogs so the observational correction can be sampled once and then reused by any statistic.

4.2 Default KNN-field completion

The default BOSS ECHOES product uses the KNN-field engine. The report’s top-level numbers are as follows. The observed catalog has 109,636 CMASS-South galaxies. The real missing-target list contains 5,272 collision targets and 1,505 redshift-failure targets. The fixed base catalog therefore has 116,413 galaxies. With the imaging-systematic analog sampler enabled, seed 0 has 119,923 galaxies, including 3,510 analogs. Seeds 0–4 vary only at the $\simeq 0.1\%$ level in total count, as expected for the smooth analog step.

Figure 3 shows the restored-target population and sampled $n(z)$. Figure 4 compares angular and two-dimensional clustering. The completed equal-weight catalogs reproduce the official weighted angular clustering to roughly 1.5% over the plotted angular range and the projected and multipole statistics at the 1–5% level, depending on scale and statistic. This is the expected accuracy target: the completed catalog is not an exact algebraic rewrite of the official weights, because it assigns redshifts and positions to missing galaxies rather than moving all their weight onto observed neighbors, but it should remain inside the systematic scale already tolerated by the reference BOSS treatment.

4.3 Truth-known validation

Two truth-known tests are used. The first hides the spectroscopic redshifts of a subset of real BOSS galaxies and then asks the completion machinery to recover the clustering of the original full catalog. The second applies the same machinery to independent MultiDark-Patchy mocks. In both tests, an “oracle” catalog that restores the hidden objects with their true redshifts provides a floor: if the oracle differs from truth, the difference comes from the masking and bookkeeping rather than from redshift assignment.

The main result is shown in Figure 5. The real-BOSS-truth mock and the Patchy mocks recover $w_p(r_p)$ to about 1–2% over $0.5 < r_p < 40 h^{-1}\text{Mpc}$. The oracle ratio lies in the range 0.997–1.005, so the remaining error is dominated by the redshift assignment posterior. The Patchy test is faithful above about $1 h^{-1}\text{Mpc}$; the sub-Mpc deficit is a known one-halo limitation of the Patchy mocks rather than a failure mode inferred from real BOSS data.




Figure 4: Two-point validation on the real BOSS CMASS-South data. The equal-weight completed catalogs reproduce the official weighted angular and redshift-space clustering while replacing completeness weights by explicit galaxies.




Figure 5: Truth-known recovery tests. The completed ensemble recovers the input projected clustering at the 1–2% level over the validated scales. The oracle curve isolates the part of the residual that comes from redshift assignment.




Figure 6: Higher-order validation with nearest-neighbor CDFs and counts-in-cells. These diagnostics probe nonlinear responses to duplicate points, missing close pairs, and redshift smearing.

4.4 Higher-order and non-two-point checks

The validation intentionally includes summaries that are nonlinear in the point set. For nearest-neighbor CDFs, the maximum absolute CDF differences between completed and truth catalogs are 0.0022, 0.0021, and 0.0011 for $k = 1, 2, 4$, respectively. For counts in $8 h^{-1}\text{Mpc}$ cells, the mean is 0.660 versus 0.663 in truth, the variance-to-mean ratio is 2.37 versus 2.42, and the skewness is 2.90 versus 2.94. These tests are not exhaustive, but they are important because a product that only reproduces w_p could still fail for nearest-neighbor or cell-count statistics.

4.5 Calibration of completion uncertainty

The completion posterior is intentionally narrower than a full survey posterior. It conditions on the observed CMASS-South realization and resamples only the unobserved spectroscopic information and smooth imaging analogs. In the Patchy calibration suite, the posterior uncertainty in $w_p(r_p)$ is 0.17–0.44% per bin, while sample variance over the same bins is 0.8–7.2%. The completion therefore adds only 6–36% of the sample-variance error in quadrature, depending on scale. The 68% completion band brackets the mock truth in only about 8% of bins because it is not intended to regenerate the large-scale density field. This is correct behavior: a user should add the completion covariance to a cosmic-variance or simulation covariance, not replace that covariance with the completion ensemble.

4.6 KNN-field versus graphGP

The KNN-field and graphGP routes agree on the main scientific conclusion but make different compromises. The KNN-field engine is sharper on the sub-Mpc collision scale because it directly uses nearby observed redshifts and the empiric close-pair prior. graphGP is smoother and more explicitly correlated across missing objects because all lines of sight share a density-field draw. In truth recovery, both routes recover the input clustering to the




Figure 7: Calibration of the completion uncertainty relative to mock sample variance. The ensemble spread is the uncertainty in the observational completion conditioned on the observed catalog, not the total survey covariance.

percent level. In the real-data head-to-head comparison, graphGP lies closer to the official weighted reference on larger projected scales, with a typical ratio near unity, while the KNN-field realization is modestly high in some bins. Because the KNN-field route is much cheaper and compresses to the 2 MB public ECHOES posterior, it is the default release. graphGP is retained as a first-class alternative for surveys or statistics where correlated redshift-completion uncertainty is essential.

5 Sampling the Completion Posterior

For the default ECHOES release, a posterior sample is a completed catalog seed. The mean of a statistic over seeds estimates the statistic for the completed survey, and the covariance over seeds estimates uncertainty due to the completion. If $S[\mathcal{C}^{(s)}]$ is any statistic of a catalog, then

$$\bar{S} = \frac{1}{N_s} \sum_{s=1}^{N_s} S[\mathcal{C}^{(s)}], \quad (14)$$

$$C_{ab}^{\text{comp}} = \frac{1}{N_s - 1} \sum_s \left(S_a[\mathcal{C}^{(s)}] - \bar{S}_a \right) \left(S_b[\mathcal{C}^{(s)}] - \bar{S}_b \right). \quad (15)$$

This covariance should be added to, or modeled jointly with, the ordinary survey covariance. It should not be interpreted as the covariance of the cosmological density field.

For graphGP the same equations apply, but the meaning of a seed changes. A seed selects a conditional density-field draw as well as the missing-galaxy redshift draws from Equation (13). The redshift assignments are therefore correlated across missing galaxies. This is the more faithful posterior object if the statistic is sensitive to coherent redshift uncertainty, but it is more expensive to construct and store. For the BOSS CMASS-South tests in this paper, the difference between graphGP and KNN-field is smaller than the sample-variance error budget for the standard clustering statistics.

The report evaluates several statistic-level responses. The angular $w(\theta)$ response to shifting the missing-galaxy





Figure 8: Comparison of the default KNN-field redshift engine with the graphGP conditional-field route. The two methods are complementary: KNN-field is fast and local; graphGP is correlated through a shared posterior density-field draw.

redshift prior is far below the realization scatter, as expected for a mostly angular statistic. Radial statistics are more sensitive. The completed $w_p(r_p)$ correction uncertainty is at the 0.2–0.4% level, and the real-data comparison with official weighted BOSS clustering is at the 1–5% level across the displayed projected and multipole statistics. For nonlinear summaries, the nearest-neighbor CDF and counts-in-cells tests in Figure 6 show no large hidden failure from the local analogs or from the missing-redshift posterior. Users can repeat the same exercise for any proposed summary by running that summary on seeds $0, \dots, N_s - 1$.

6 Systematics, Masks, and Trustworthiness

The completion also produces diagnostic maps rather than only final catalogs. The redshift-failure coupling to the local target density is consistent with zero in the current BOSS CMASS-South product, $h = -0.06 \pm 0.06$, while the fiber-collision coupling is detected, $h = -0.24 \pm 0.03$. This is physically reasonable: collisions occur in denser angular environments. The completed density map correlates with the total-target density at 0.98, and the trustworthiness map has a median scatter of 0.13. These diagnostics identify where the completion is extrapolating most strongly.

The optional mask-inpainting product fills 542 interior holes totaling 16.6 deg^2 , about 0.9% of the footprint. The inpainting closure test is approximately 1% over 0.06° – 2° . We keep this step optional because different analyses may prefer to leave bright-star holes and small mask features explicit in the random catalog. The default product therefore corrects spectroscopic incompleteness and smooth imaging weights while leaving the official footprint visible.



7 Data and Code Availability

The analysis code and ECHOES report products are in the public repository <https://github.com/yipihey/graphGP-cosmology>. The BOSS completion report used for this manuscript is <https://yipihey.github.io/graphGP-cosmology/completion.html>. The compact posterior and sampler are distributed in the report’s docs/data/ directory:

```
cmass_south_posterior.npz
cmass_south_randoms.npz
draw_samples.py
```

The minimal use pattern is

```
from draw_samples import load_package, draw
pkg = load_package("cmass_south_posterior.npz")
cat = draw(pkg, seed=0)
```

which returns arrays `ra`, `dec`, `z`, `prov`, and `N`. A reproducible ensemble is specified by the integer seeds used.

The relevant source modules are `twopt_density/boss.py` for the BOSS loader, `twopt_density/cmass_targets.py` for missing target recovery, `twopt_density/photoz.py` for the color-space k NN photo- z , `twopt_density/observed_ls.py` for the KNN-field completion and graphGP redshift interface, `twopt_density/posterior_sampler.py` for the compact inverse-CDF posterior, `twopt_density/density_field.py` for graphGP posterior field sampling, and `twopt_density/clustering_corrfunc.py` for the standard projected and multipole clustering measurements.

8 Conclusions

We have presented ECHOES, an equal-weight completed-catalog posterior for BOSS DR12 CMASS-




Figure 9: Systematics diagnostics for the completion. The maps and coupling tests quantify where the data-driven completion is well anchored by observed targets and where it is most dependent on the model.

South. The product restores spectroscopically missing galaxies at real SDSS imaging positions, samples their redshifts from a data-driven posterior, and represents smooth imaging-systematic weights as local analog galaxies. It therefore changes the user-facing data interface: instead of asking every new summary statistic to understand the BOSS completeness weights, an analyst can run the statistic directly on completed point catalogs and propagate observational-completion uncertainty by resampling seeds.

The validation is deliberately conservative. The completed ensemble recovers truth-known $w_p(r_p)$ to the 1–2% level over the validated scales, matches official weighted BOSS clustering at the few-percent level, reproduces nearest-neighbor CDFs and counts-in-cells diagnostics, and has a completion uncertainty smaller than the CMASS-South sample variance. The graphGP route provides an independent, correlated posterior construction and agrees with the default KNN-field route within the current BOSS error budget.

The broader implication is practical. Survey collaborations have learned how to publish two-point-ready data products. The next generation of analyses will not be limited to two-point functions, and the next generation of surveys will have more complicated assignment, calibration, and selection functions. A ECHOES posterior over completed survey catalogs is a way to preserve the historical division of labor—survey teams ground observational corrections in the data, theory teams test statistics and models—while making the interface useful for arbitrary summaries. DESI, Euclid, Roman, CSST, and Rubin/LSST will all face versions of this problem. BOSS CMASS-South is small enough to solve carefully and mature enough to validate against the standard answer; it is therefore a useful prototype for how future public survey products can serve both standard and beyond-two-point cosmology.

A Reproducible Algorithm

This appendix writes the default KNN-field completion as an implementation recipe.

1. Load the BOSS DR12 CMASS-South LSS catalog with photometry and the matching random catalog. Apply the cuts $RA < 28^\circ$ or $RA > 335^\circ$, $Dec > -6^\circ$, and $0.45 < z < 0.60$.

2. Convert SDSS MODELFLUX and EXTINCTION to magnitudes and colors. Train PhotoZKNN($k=100$) on finite observed galaxies using features $(g-r, r-i, i-z, i)$.
3. Load the CMASS target catalog. Build the missing list by greedily tying no-spectrum targets to WEIGHT_CP hosts within $62''$ and ZWARNING $\neq 0$ targets to WEIGHT_NOZ hosts within $15'$.
4. Measure the observed close-pair Δz pool from pairs of observed galaxies within $62''$ and symmetrize it.
5. For each missing target, find the $K = 150$ angular nearest observed spectroscopic galaxies. On a 256-point redshift grid, form the product of the local field KDE with $\sigma_f = 0.004$, the photo- z KDE with $\sigma_p = 0.02$, and, for collisions, the close-pair prior.
6. Store each normalized posterior as a 65-level inverse CDF. Quantize the observed redshifts and inverse-CDF values to `uint16`; store positions as `float32`, WEIGHT_SYSTOT as `float16`, and provenance as `int8`.
7. To draw seed s , invert each missing target’s CDF using uniform variates from `numpy.random.default_rng(s)`, add Gaussian redshift jitter $\sigma_f/2$, concatenate observed galaxies and restored targets, and add local imaging-systematic analogs with $1''$ angular jitter.
8. Measure statistics on many seeds. Use the seed mean as the completed survey statistic and the seed covariance as the completion covariance.

B graphGP Implementation Details

The graphGP route uses the same observed BOSS catalog and missing target list but replaces the KNN line-of-sight density with a correlated posterior field. The implementation steps are:

1. Choose a fiducial distance-redshift relation for the internal field coordinate system and convert observed (α, δ, z) to $\mathbf{x}$.
2. Measure or tabulate the covariance $K_\xi(r)$ from the observed two-point function.

3. Build a sparse Vecchia neighbor graph over data and prediction points.
4. Generate a GP prior draw on that graph and condition it with Equation (12) using the Gaussianized Poisson likelihood.
5. Store $1 + \delta_{\text{post}}$ on a redshift-shell by HealPIX light cone.
6. For each missing target, evaluate Equation (13) on the same redshift grid used by the KNN-field sampler, then draw redshifts from the shared field realization.

The field posterior is not part of the 2 MB default release because it is larger and more expensive, but it is implemented as a first-class alternative and was used for the comparisons in Section 4.6.

C Suggested Figure Set

The current report contains more validation figures than should appear in the main paper. The main text uses the footprint/weights, photo- z , missing targets, two-point validation, truth recovery, higher-order validation, calibration, graphGP comparison, and systematics figures. A shorter journal version could move Figure 9 and the optional inpainting diagnostics to an online appendix. The figures `mask.png`, `inpaint_closure.png`, `cosmology_consistency.png`, and `graphgp_cosmology_invariance.png` should be retained as supplementary material because they answer likely referee questions about mask holes and the fiducial cosmology used internally by graphGP.

Acknowledgments

This manuscript draft is based on the BOSS completion products in the `graphGP-cosmology` repository. The BOSS data products are from SDSS-III.

Software

Corrfunc [Sinha and Garrison, 2020], NumPy, SciPy, Astropy, Healpy, graphGP

References

- S. Alam et al. The clustering of galaxies in the completed sdss-iii baryon oscillation spectroscopic survey: cosmological analysis of the dr12 galaxy sample. *Monthly Notices of the Royal Astronomical Society*, 470:2617, 2017.
- C. Alcock and B. Paczynski. An evolution free test for non-zero cosmological constant. *Nature*, 281:358, 1979.
- A. Banerjee and T. Abel. Nearest neighbour distributions: New statistical measures for cosmological clustering. *Monthly Notices of the Royal Astronomical Society*, 500:5479, 2021.
- C. M. Baugh, E. Gaztanaga, and G. Efstathiou. A comparison of the evolution of density fields in perturbation theory and numerical simulations. ii. counts-in-cells analysis. *Monthly Notices of the Royal Astronomical Society*, 274:1049, 1995.
- Beyond-2pt Collaboration, E. Krause, et al. A parameter-masked mock data challenge for beyond-two-point galaxy clustering statistics. *The Astrophysical Journal*, 990:99, 2025.
- M. Biagetti, A. Cole, and G. Shiu. Fisher forecasts for primordial non-gaussianity from persistent homology. *Journal of Cosmology and Astroparticle Physics*, 2022:002, 2022.
- S. Cheng, Y.-S. Ting, B. Ménard, and J. Bruna. A new approach to observational cosmology using the scattering transform. *Monthly Notices of the Royal Astronomical Society*, 499:5902, 2020.
- S. Cole et al. The 2df galaxy redshift survey: power-spectrum analysis of the final data set and cosmological implications. *Monthly Notices of the Royal Astronomical Society*, 362:505, 2005.
- M. Colless et al. The 2df galaxy redshift survey: spectra and redshifts. *Monthly Notices of the Royal Astronomical Society*, 328:1039, 2001.
- CSST Collaboration, Y. Gong, et al. Introduction to the chinese space station survey telescope. *Science China Physics, Mechanics, and Astronomy*, 69:39501, 2026.
- M. Davis and P. J. E. Peebles. A survey of galaxy redshifts. v. the two-point position and velocity correlations. *The Astrophysical Journal*, 267:465, 1983. doi: 10.1086/160884.
- K. S. Dawson et al. The baryon oscillation spectroscopic survey of sdss-iii. *The Astronomical Journal*, 145:10, 2013. doi: 10.1088/0004-6256/145/1/10.
- V. de Lapparent, M. J. Geller, and J. P. Huchra. A slice of the universe. *The Astrophysical Journal Letters*, 302:L1, 1986. doi: 10.1086/184625.
- M. Einasto, J. Einasto, E. Tago, G. B. Dalton, and H. Andernach. Structure and formation of superclusters. xiii. the void probability function. *Monthly Notices of the Royal Astronomical Society*, 248:593, 1991.
- D. J. Eisenstein et al. Detection of the baryon acoustic peak in the large-scale correlation function of sdss luminous red galaxies. *The Astrophysical Journal*, 633:560, 2005. doi: 10.1086/466512.

- Euclid Collaboration, Y. Mellier, et al. Euclid: I. overview of the euclid mission. *Astronomy and Astrophysics*, 697:A1, 2025.
- H. A. Feldman, N. Kaiser, and J. A. Peacock. Power-spectrum analysis of three-dimensional redshift surveys. *The Astrophysical Journal*, 426:23, 1994.
- O. Friedrich, C. Uhlemann, F. Villaescusa-Navarro, et al. Primordial non-gaussianity without tails – how to measure fnl with the bulk of the density pdf. *Monthly Notices of the Royal Astronomical Society*, 498:464, 2020.
- J. N. Fry and E. Gaztanaga. Biasing and hierarchical statistics in large-scale structure. *The Astrophysical Journal*, 413:447, 1993.
- H. Guo, I. Zehavi, and Z. Zheng. A new method to correct for fiber collisions in galaxy two-point statistics. *The Astrophysical Journal*, 756:127, 2012.
- C. Hahn et al. Simbig: mock challenge for a forward modeling approach to galaxy clustering. *Journal of Cosmology and Astroparticle Physics*, 2023:010, 2023a.
- C. Hahn et al. A forward modeling approach to analyzing galaxy clustering with simbig. *Proceedings of the National Academy of Sciences*, 120:e2218810120, 2023b.
- C. Hahn et al. Cosmological constraints from the nonlinear galaxy bispectrum. *Physical Review D*, 109:083534, 2024a.
- C. Hahn et al. Cosmological constraints from non-gaussian and nonlinear galaxy clustering using the simbig inference framework. *Nature Astronomy*, 8:1457, 2024b.
- A. J. S. Hamilton. Toward better ways to measure the galaxy correlation function. *The Astrophysical Journal*, 417:19, 1993.
- G. Harker, S. Cole, J. Helly, C. Frenk, and A. Jenkins. A marked correlation function analysis of halo formation times in the millennium simulation. *Monthly Notices of the Royal Astronomical Society*, 367:1039, 2006.
- Edwin Hubble. The distribution of extra-galactic nebulae. *The Astrophysical Journal*, 79:8, 1934. doi: 10.1086/143517.
- J. Jasche and B. D. Wandelt. Bayesian physical reconstruction of initial conditions from large-scale structure surveys. *Monthly Notices of the Royal Astronomical Society*, 432:894, 2013.
- N. Kaiser. Clustering in real space and in redshift space. *Monthly Notices of the Royal Astronomical Society*, 227:1, 1987.
- M. Katzfuss, J. Guinness, W. Gong, and D. Zilber. Vecchia approximations of gaussian-process predictions. *Journal of Agricultural, Biological and Environmental Statistics*, 25:383, 2020.
- F.-S. Kitaura et al. The clustering of galaxies in the sdss-iii baryon oscillation spectroscopic survey: mock galaxy catalogues for the boss final data release. *Monthly Notices of the Royal Astronomical Society*, 456:4156, 2016.
- T.-W. Lan et al. The desi survey validation: Results from visual inspection of bright galaxies, luminous red galaxies, and emission-line galaxies. *The Astrophysical Journal*, 943:68, 2023.
- S. D. Landy and A. S. Szalay. Bias and variance of angular correlation functions. *The Astrophysical Journal*, 412:64, 1993. doi: 10.1086/172900.
- J. Lasker et al. Production of alternate realizations of desi fiber assignment for unbiased clustering measurement in data and simulations. *Journal of Cosmology and Astroparticle Physics*, 2025:127, 2025.
- P. Lemos et al. Field-level simulation-based inference of galaxy clustering with convolutional neural networks. *Physical Review D*, 109:083536, 2024.
- E. Massara, F. Villaescusa-Navarro, S. Ho, N. Dalal, and D. N. Spergel. Using the marked power spectrum to detect the signature of neutrinos in large-scale structure. *Physical Review Letters*, 126:011301, 2021.
- E. Massara, F. Villaescusa-Navarro, C. Hahn, et al. Cosmological information in the marked power spectrum of the galaxy field. *The Astrophysical Journal*, 951:70, 2023.
- E. Massara et al. Cosmological constraints using simulation-based inference of galaxy clustering with marked power spectra. *Physical Review D*, 112:083507, 2025.
- H. Miao, Y. Gong, X. Chen, et al. Cosmological constraint precision of photometric and spectroscopic multi-probe surveys of china space station telescope. *Monthly Notices of the Royal Astronomical Society*, 519:1132, 2023.
- F. G. Mohammad et al. The completed sdss-iv extended baryon oscillation spectroscopic survey: pairwise-inverse probability and angular correction for fibre collisions in clustering measurements. *Monthly Notices of the Royal Astronomical Society*, 498:128, 2020.
- N.-M. Nguyen, F. Schmidt, G. Lavaux, and J. Jasche. How much information can be extracted from galaxy clustering at the field level? *Physical Review Letters*, 133:221006, 2024.

- L. Pinol, R. N. Cahn, N. Hand, U. Seljak, and M. White. Imprint of desi fiber assignment on the anisotropic power spectrum of emission line galaxies. *Journal of Cosmology and Astroparticle Physics*, 2017:008, 2017.
- B. Regalado-Saint Blancard et al. Galaxy clustering analysis with simbig and the wavelet scattering transform. *Physical Review D*, 109:083535, 2024.
- B. Reid et al. Sdss-iii baryon oscillation spectroscopic survey data release 12: galaxy target selection and large-scale structure catalogues. *Monthly Notices of the Royal Astronomical Society*, 455:1553, 2016.
- A. J. Ross et al. The clustering of galaxies in the completed sdss-iii baryon oscillation spectroscopic survey: observational systematics and baryon acoustic oscillations in the correlation function. *Monthly Notices of the Royal Astronomical Society*, 464:1168, 2017.
- G. B. Rybicki and W. H. Press. Interpolation, realization, and reconstruction of noisy, irregularly sampled data. *The Astrophysical Journal*, 398:169, 1992.
- S. Satpathy et al. Measurement of marked correlation functions in sdss-iii baryon oscillation spectroscopic survey using lowz galaxies in data release 12. *Monthly Notices of the Royal Astronomical Society*, 484:2148, 2019.
- F. Schäfer, M. Katzfuss, and H. Owhadi. Sparse cholesky factorization by kullback-leibler minimization. *SIAM Journal on Scientific Computing*, 43:A2019, 2021.
- D. J. Schlegel, D. P. Finkbeiner, and M. Davis. Maps of dust infrared emission for use in estimation of reddening and cosmic microwave background radiation foregrounds. *The Astrophysical Journal*, 500:525, 1998. doi: 10.1086/305772.
- J. H. Silber et al. The robotic multiobject focal plane system of the dark energy spectroscopic instrument. *The Astronomical Journal*, 165:9, 2023.
- M. Sinha and L. H. Garrison. Corrfunc – a suite of blazing fast correlation functions on the cpu. *Monthly Notices of the Royal Astronomical Society*, 491:3022, 2020.
- S. A. Smee et al. The multi-object, fiber-fed spectrographs for the sloan digital sky survey and the baryon oscillation spectroscopic survey. *The Astronomical Journal*, 146:32, 2013. doi: 10.1088/0004-6256/146/2/32.
- I. Szapudi and A. S. Szalay. Higher order statistics of the galaxy distribution using generating functions. *The Astrophysical Journal*, 408:43, 1993.
- I. Szapudi, A. Meiksin, and R. C. Nichol. Higher order statistics from the edinburgh/durham southern galaxy catalogue survey. i. counts in cells. *The Astrophysical Journal*, 473:15, 1996.
- C. Uhlemann, O. Friedrich, F. Villaescusa-Navarro, A. Banerjee, and S. Codis. Fisher for complements: extracting cosmology and neutrino mass from the counts-in-cells pdf. *Monthly Notices of the Royal Astronomical Society*, 495:4006, 2020.
- G. Valogiannis and C. Dvorkin. Towards an optimal estimation of cosmological parameters with the wavelet scattering transform. *Physical Review D*, 105:103534, 2022.
- F. Villaescusa-Navarro et al. The quijote simulations. *The Astrophysical Journal Supplement Series*, 250:2, 2020.
- Y. Wang et al. The high latitude spectroscopic survey on the nancy grace roman space telescope. *The Astrophysical Journal*, 928:1, 2022.
- G. Wilding, K. Nevenzeel, R. van de Weygaert, et al. Persistent homology of the cosmic web. i. hierarchical topology in Λ cdm cosmologies. *Monthly Notices of the Royal Astronomical Society*, 507:2968, 2021.
- J. H. T. Yip, M. Biagetti, and A. Cole. Cosmology with persistent homology: a fisher forecast. *Journal of Cosmology and Astroparticle Physics*, 2024:034, 2024.
- D. G. York et al. The sloan digital sky survey: Technical summary. *The Astronomical Journal*, 120:1579, 2000. doi: 10.1086/301513.
- S. Yuan, T. Abel, and R. H. Wechsler. Robust cosmological inference from non-gaussian and nonlinear galaxy clustering using kth-nearest-neighbour statistics. *Monthly Notices of the Royal Astronomical Society*, 527:1993, 2024.
- H. Zhan and J. A. Tyson. Cosmology with the large synoptic survey telescope: an overview. *Reports on Progress in Physics*, 81:066901, 2018.